

# Computer Vision: Lab 1

This lab includes three parts. We will use the Matlab environment throughout. Part one uses different kinds of filters to manipulate the image. In part two, we will write a convolution function to understand the process of convolution. Part three introduces an example of fitting a line to detected edges. Attached files are two example images, and the skeleton `lab1part3.m`. You may wish to read the documentation on matlab functions mentioned in the following by `doc functionname`. The instructions also include supplementary exercises for keen students. These can be skipped first, and tried after the core exercises are complete.

## Part 1: Warmup with Convolution and Correlation

### 1.1 Denoising

In this section, we will warm up by trying different filters to manipulate the image. You can read an image into matlab, and view it using commands like:

```
img = double(rgb2gray(imread('picture/input.png')))/255;
imshow(img);
```

We can create a noisy version of this image to work with denoising:

```
imgn = imnoise(img, 'salt & pepper', 0.02);
imshow(imgn);
```

Now we can make a 3x3 box filter and convolve it with the image for denoising.

```
filter = 1/9*[1 1 1;1 1 1;1 1 1];
imgdn = conv2(imgn, filter, 'same');
imshow([imgn, imgdn]);
```

You should see a denoised but slightly blurred version of the input image.

Points to Explore:

1. Increase the size of the box filter (5x5, 7x7, etc) and observe the tradeoff between increased denoising effect and increased blurring effect.
  - Hint 1: You can use `ones(h,w)` to generate a size `h`x`w` matrix of 1s.
  - Hint 2: Remember to normalise the filter so that the values sum to 1.0.
2. Verify that you can achieve the same effects with the matlab correlation rather than convolution routine `filter2()`.
  - Hint 1: `filter2` expects image/filter arguments in reverse order to `conv2`.
  - Hint 2: You can compare the numeric similarities of two matrices like: `norm(imgdn(:)-imgdn2(:))`. The `(:)` notation reads out a matrix as a vector, and `norm()` gets the magnitude of a vector.
  - Why do convolution and correlation give the same result here?
3. Compare the output of convolution with a Gaussian kernel to the box filter kernel.

- Hint: You can use `filter=fspecial('gaussian',h,s)` to create a Gaussian filter. You can also visualise the filter with `imshow(filter)`.
  - Observe that changing the standard deviation of the Gaussian can also be used to create different degrees of blur.
    - Bonus [Moderate] : Implement your own Gaussian kernel rather than using `fspecial`.
4. Use the median filtering routine `medfilt2` to achieve more effective de-noising with less blurring.


## 1.2 Edge Detection

Now define and apply an edge detection filter like:

```
img =double(rgb2gray(imread('picture/input.png')))/255;
filter_edge = [-1,0,1];
img_edge = conv2(img, filter_edge, 'same');
imshow([img,img_edge])
```

Points to Explore:

1. Modify this example to perform horizontal edge detection.
2. Filter (`filter2`) rather than convolve the image with `filter_edge`.
  - This time the output of `filter2` and `conv2` are different. Why?
  - What has `filter2` computed instead?
  - Adjust the call to `filter2` to make it produce exactly the same output as `conv2`.
  - Hint: Check out `rot90()`.
3. Compute the overall magnitude of edge pixels in both x and y directions. Recall the formula:  $\|\nabla f\| = \sqrt{(\frac{\partial f}{\partial x})^2 + (\frac{\partial f}{\partial y})^2}$ . It should look like this, showing the original image, the vertical edges, the horizontal edges, and the magnitude (the edges are multiplied by 3.0 to increase visual contrast):



4. Explore the impact of filtering the image first with a blur filter (box or Gaussian), before performing edge detection, as discussed in the lecture notes. You should be able to use this to reduce the “noise” edges, and detect thicker/thinner edges.
5. Bonus [Moderate]: Verify the associativity of convolution numerically. If you convolve the edge filter and blur filter first, then convolve the result with the image; it should achieve the same result as sequentially blurring then edge detecting. You can use `tic` and `toc` in matlab to measure execution time. Which way is faster?

6. Bonus [Moderate]: Compute and visualise the edge direction at each pixel.

## Part 2: Understanding Convolution

To make sure you understand convolution properly, implement your own convolution routine that inputs an image, a filter, and returns an image that is the result of convolving the image with the filter.

- For the simplest implementation, it is probably easiest to use the semantics of `shape='valid'` (see lecture notes, documentation of matlab `conv2`) so the output image is the difference of the input image and filter size.
- To check your implementation, take an example image and filter, and verify that your routine gives the same output as the built in matlab `conv2`.
  - a. Hint: Call `conv2` with the same shape parameter that matches the edge semantics of your convolution routine.
- Bonus [Advanced]: Vanilla convolution takes  $O(N^2M^2)$  cost. Implement fast  $O(N\log(N)M\log(M))$  convolution with FFT (Hint: matlab `fft2`)  
[https://en.wikipedia.org/wiki/Convolution\\_theorem](https://en.wikipedia.org/wiki/Convolution_theorem)

## Part 3: Line Detection

In this section we will detect edges in a road image and use these to fit lines to the image. A self-driving car may need to do this in order to localise the road for steering. Open the provided skeleton code `lab1part3.m`. This code performs edge detection, and extracts a list of coordinates of every edge point in the left (`edge1`) and right (`edge2`) parts of the image. These will correspond to the two sides of the road.

- Implement the simple least squares line-fitting from the lecture notes to fit a line to the left and right sides of the road. It should look like this:



- Bonus [Advanced]: Implement a RANSAC line detector for robust line fitting.